


Vodis

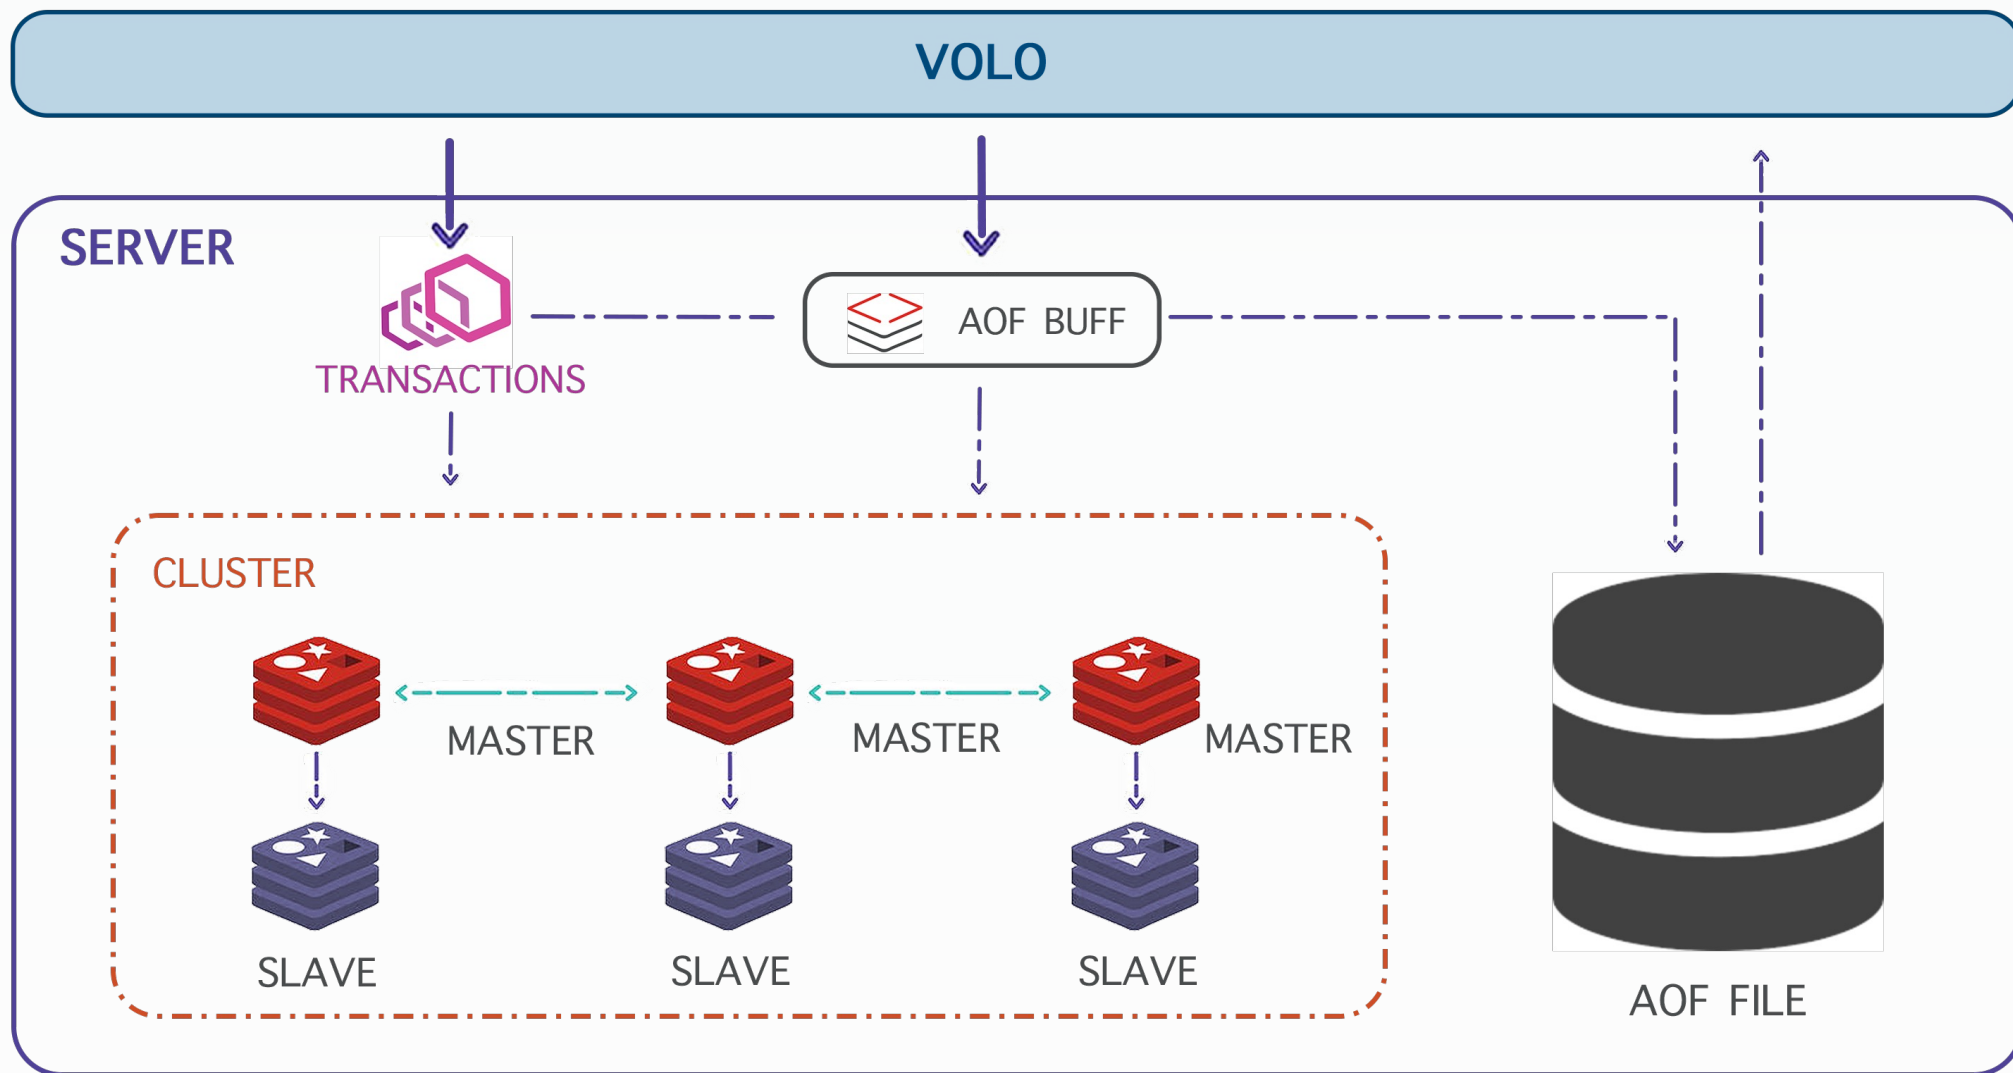
基于VOLO-RPC的Mini-Redis实现



Vodis 架构



浙江大学
ZHEJIANG UNIVERSITY





文件结构

- > **idl**
- > **src**
- > **statics**
- > **volo-gen**
- . gitignore
- Cargo.lock
- Cargo.toml
- README.md
- clear.sh
- cluster.toml
- get.sh
- master_slave_test.sh
- rust-toolchain.toml
- set.sh
- simple_transaction.sh
- start_main_server.sh
- start_slave_server.sh
- test1.sh
- transaction_with_watch.sh

- ▼ **idl**
 - vol_redis.thrift

- ▼ **src**
 - > **bin**
 - cmdargs.rs Viewer
 - lib.rs Controller
 - redis.rs Model

- ▼ **bin**
 - client-cli.rs
 - proxy.rs

```
async fn subscribe(handle: String) -> ! {
```

```
pub fn get_client(addr: impl Into<Address>) -> volo  
    volo_gen::volo::redis::ItemServiceClientBuilder
```

```
tokio::spawn(future: async move {  
    //AOF线程, 实现everysec策略
```

1.AOF (Append-only File) 实现持久化 (30分)

1.对于一个已经存在的 log file (格式自定, 文件自行准备), 在 redis server 启动的时候可以根据该文件重建redis 数据。

2.提供一个测试脚本:

Client 发送一些 Set 操作之后, 重启 Server, 然后再用 client 进行 Get 调用, 可以访问到之前的数据

```
vodis[connected]> set 1 12312 ex 2
[OK] OK
vodis[connected]> set 2 213132
[OK] OK
vodis[connected]> set 3 11451
[OK] OK
vodis[connected]> set 4 1233 ex 100000
[OK] OK
vodis[connected]> del 2
[OK] 1
(deleted count)
vodis[connected]>
```

```
^Z
[1] + 4768 suspended cargo run --bin server -- -i 127.0
→ volo-mini-redis git:(main) X kill 4768
→ volo-mini-redis git:(main) X
```

```
Warning: mini_redis (bin client-cli) generated 1 warning
Finished dev [unoptimized + debuginfo] target(s) in 4.42s
Running `target/debug/client-cli -s '127.0.0.1:8080'`
vodis[connected]> get 1
[FAILED] (nil)
vodis[connected]> get 2
[FAILED] (nil)
vodis[connected]> get 3
[OK] 11451
vodis[connected]> get 4
[OK] 1233
vodis[connected]>
```

1.AOF (Append-only File) 实现持久化 (30分)

亮点:

1.AOF支持Expire time;

2.AOF支持两种写盘策略:

混合模式(低负载) ——定时(1s)刷盘

高负载——缓冲区满刷盘

```
vodis[connected]> set 1 12312 ex 2
[OK] OK
vodis[connected]> set 2 213132
[OK] OK
vodis[connected]> set 3 11451
[OK] OK
vodis[connected]> set 4 1233 ex 100000
[OK] OK
vodis[connected]> del 2
[OK] 1
(deleted count)
vodis[connected]>
```

```
^Z
[1] + 4768 suspended cargo run --bin server -- -i 127.0
→ volo-mini-redis git:(main) X kill 4768
→ volo-mini-redis git:(main) X
```

```
Warning: mini_redis (bin client-cli) generated 1 warning
Finished dev [unoptimized + debuginfo] target(s) in 4.42s
Running `target/debug/client-cli -s '127.0.0.1:8080'`
vodis[connected]> get 1
[FAILED] (nil)
vodis[connected]> get 2
[FAILED] (nil)
vodis[connected]> get 3
[OK] 11451
vodis[connected]> get 4
[OK] 1233
vodis[connected]>
```

2.Redis 主从架构 (30分)

- 1.定义一个配置文件的格式，可以从该配置文件启动多个 redis 实例，并且区分哪些是主节点，哪些是从节点。
- 2.从节点对于 Set 操作返回错误。
- 3.对主节点进行 Set 操作，可以通过从节点通过 Get 获取到设置的数据。

亮点:

- 1.支持一主多从，从上加从
- 2.支持集群间事务(通过主为从分发的uuid实现指令来源区分)

```
warning: `mini_redis` (bin "client-cli") generated 1 warning (run `cargo fix`  
Finished dev [unoptimized + debuginfo] target(s) in 2.76s  
Running `target\debug\client-cli.exe -s 127.0.0.1:8080`  
vredis[connected]> get 1  
[OK] 4617238123  
vredis[connected]> set 1 12312312  
[OK] OK  
vredis[connected]> 
```

```
warning: `mini_redis` (bin "client-cli") generated 1 warning (run `cargo fix`  
Finished dev [unoptimized + debuginfo] target(s) in 0.32s  
Running `target\debug\client-cli.exe -s 127.0.0.1:8888`  
vredis[connected]> get 1  
[OK] 4617238123  
vredis[connected]> get 1  
[OK] 12312312  
vredis[connected]> 
```

```
warning: `mini_redis` (lib) generated 8 warnings (run `cargo fix --lib -p mini_redis`  
Finished dev [unoptimized + debuginfo] target(s) in 0.30s  
Running `target\debug\server.exe -i 127.0.0.1 -p 8888 --slaveof 127.0.0.1:8080`  
Sync send.  
2023-09-15T02:35:29.126315Z INFO volo_thrift::server: [VOLO] server start at: Tcp(Tc  
0.0.1:8888, socket: 344 }) } })  
2023-09-15T02:35:29.714687Z INFO mini_redis: Serialize success!!!  
xxx  
COMMAND SAVED!!  
propagate to 0 clients  
xxx  
COMMAND SAVED!!  
propagate to 0 clients  
 
```

3.Redis Cluster (40分)

- 1.实现一个 Redis Proxy, 可以将 Set 和 Get 请求转发到其他的 Redis 实例上进行处理。
- 2.定义一个配置文件的格式, 可以从该配置文件启动多个 redis 实例和一个 redis proxy。
- 3.Redis Proxy 可以根据 Get 和 Set 目标的 key 将请求分发到不同的 redis 实例进行处理。



测试样例

3.Redis Cluster (40分)

8080 主1-8082从1

8081 主2-8083从2

proxy: 8888

get/set经过proxy→主→从 多层转发。

```
70 cargo run --bin proxy -- --cfg cluster.toml # 挂载到proxy服务器

问题 11 输出 调试控制台 端口 终端

| help: remove this `mut`
|
warning: variable does not need to be mutable
--> src/lib.rs:744:21
744 | ...et mut transaction = tran...
|      ^^^^^^^^^^^^^
| help: remove this `mut`
|
warning: variable does not need to be mutable
--> src/lib.rs:750:29
750 | ... let mut watch_pair = tra...
|      ^^^^^^^^^^^^^
| help: remove this `mut`
|

warning: `mini_redis` (lib) generated 8 warnings (run `cargo fix --lib -p mini_redis` to apply 8 suggestions)
Finished dev [unoptimized + debuginfo] target(s) in 0.41s
Running `target/debug/server -i 127.0.0.1 -p 8888 --name proxy`
2023-09-15T02:53:43.875671Z INFO volo::hotrestart: hot_restart skip dup_parent_listener_sock: 127.0.0.1:8888, as parent
2023-09-15T02:53:43.875885Z INFO volo::hotrestart: hot_restart register_listener_fd: 127.0.0.1:8888, 10
2023-09-15T02:53:43.875980Z INFO volo_thrift::server: [VOL0] server start at: Tcp(TcpListenerStream { inner: PollEvented { io: Some(TcpListener { addr: 127.0.0.1:8888, fd: 10 }) }) })

[OK] OK
warning: `mini_redis` (bin "proxy") generated 2 warnings (run `cargo fix --bin "proxy"` to apply 1 suggestion)
Finished dev [unoptimized + debuginfo] target(s) in 0.48s
Running `target/debug/proxy -a '127.0.0.1:8888' --masters '127.0.0.1:8080' --masters '127.0.0.1:8081'`
vredis[cluster:proxy]> set a a
2023-09-15T02:55:11.473978Z INFO proxy: proxied to 127.0.0.1:8081.
[OK] OK
vredis[cluster:proxy]> set b b
2023-09-15T02:55:14.461754Z INFO proxy: proxied to 127.0.0.1:8081.
[OK] OK
vredis[cluster:proxy]> set c c
2023-09-15T02:55:16.848329Z INFO proxy: proxied to 127.0.0.1:8081.
[OK] OK
vredis[cluster:proxy]> set d d
2023-09-15T02:55:19.449357Z INFO proxy: proxied to 127.0.0.1:8080.
[OK] OK
vredis[cluster:proxy]> set e e
2023-09-15T02:55:21.979445Z INFO proxy: proxied to 127.0.0.1:8080.
[OK] OK
vredis[cluster:proxy]> get nope
[FAILED] (nil)
2023-09-15T02:55:27.422097Z INFO proxy: (by Bloom)
vredis[cluster:proxy]> get e
2023-09-15T02:55:30.849526Z INFO proxy: proxied to 127.0.0.1:8080.
[OK] e
vredis[cluster:proxy]> get a
2023-09-15T02:55:33.921922Z INFO proxy: proxied to 127.0.0.1:8081.
[OK] a
vredis[cluster:proxy]>

[+] [-] ... ^ [-] x
[-] mas1
[-] mas2
[-] sla1
[-] sla2
[-] proxy_srv
[-] proxy_client
```



测试样例

3.Redis Cluster (40分)

8080 主1-8082从1

8081 主2-8083从2

proxy: 8888

get/set经过proxy→主→从 多层转发。

```
问题 11 输出 调试控制台 端口 终端

750 | ... let mut watch_pair = tra...
    |             ^^^^^^^^^^^^^^^^^
    |             help: remove this `mut`

warning: `mini_redis` (lib) generated
8 warnings (run `cargo fix --lib -p
mini_redis` to apply 8 suggestions)
Finished dev [unoptimized + debug
info] target(s) in 0.33s
Running `target/debug/server -i
127.0.0.1 -p 8080 --name mas1`
2023-09-15T02:51:02.483922Z INFO vol
o::hotrestart: hot_restart skip dup_p
arent_listener_sock: 127.0.0.1:8080,
as parent
2023-09-15T02:51:02.484126Z INFO vol
o::hotrestart: hot_restart register_l
istener_fd: 127.0.0.1:8080, 10
2023-09-15T02:51:02.484211Z INFO vol
o_thrift::server: [VOL0] server start
at: Tcp(TcpListenerStream { inner: P
ollEvented { io: Some(TcpListener { a
ddr: 127.0.0.1:8080, fd: 10 }) }) })
xxx
propagate to 0 clients
COMMAND SAVED!!
xxx
propagate to 0 clients
COMMAND SAVED!!
[]

127.0.0.1 -p 8081 --name mas2`
2023-09-15T02:48:58.043948Z INFO vol
o::hotrestart: hot_restart skip dup_p
arent_listener_sock: 127.0.0.1:8081,
as parent
2023-09-15T02:48:58.044113Z INFO vol
o::hotrestart: hot_restart register_l
istener_fd: 127.0.0.1:8081, 10
2023-09-15T02:48:58.044182Z INFO vol
o_thrift::server: [VOL0] server start
at: Tcp(TcpListenerStream { inner: P
ollEvented { io: Some(TcpListener { a
ddr: 127.0.0.1:8081, fd: 10 }) }) })
2023-09-15T02:51:44.749322Z ERROR vol
o_thrift::transport::pool: [VOL0] cre
ate connection error: Os { code: 61,
kind: ConnectionRefused, message: "Co
nnection refused" }, key: Ip(127.0.0.
1:8083)
xxx
propagate to 1 clients
set... to 127.0.0.1:8083.
COMMAND SAVED!!
xxx
propagate to 1 clients
set... to 127.0.0.1:8083.
COMMAND SAVED!!
xxx
propagate to 1 clients
set... to 127.0.0.1:8083.
COMMAND SAVED!!
[]

warning: variable does not need to be
mutable
--> src/lib.rs:750:29

750 | ... let mut watch_pair = tra...
    |             ^^^^^^^^^^^^^^^^^
    |             help: remove this `mut`

warning: `mini_redis` (lib) generated
8 warnings (run `cargo fix --lib -p
mini_redis` to apply 8 suggestions)
Finished dev [unoptimized + debug
info] target(s) in 0.95s
Running `target/debug/server -s
'127.0.0.1:8081' --name sla1 -i 127.0
.0.1 -p 8082`
Sync send.
2023-09-15T03:01:07.734632Z INFO vol
o::hotrestart: hot_restart skip dup_p
arent_listener_sock: 127.0.0.1:8082,
as parent
2023-09-15T03:01:07.734781Z INFO vol
o::hotrestart: hot_restart register_l
istener_fd: 127.0.0.1:8082, 10
2023-09-15T03:01:07.734841Z INFO vol
o_thrift::server: [VOL0] server start
at: Tcp(TcpListenerStream { inner: P
ollEvented { io: Some(TcpListener { a
ddr: 127.0.0.1:8082, fd: 10 }) }) })
[]

p mini_redis` to apply 8 suggestions
)
Finished dev [unoptimized + debu
ginfo] target(s) in 0.36s
Running `target/debug/server -i
127.0.0.1 -p 8083 --slaveof '127.0.
0.1:8081' --name sla2`
Sync send.
2023-09-15T02:51:44.749960Z INFO vo
lo::hotrestart: hot_restart skip dup
_parent_listener_sock: 127.0.0.1:808
3, as parent
2023-09-15T02:51:44.750086Z INFO vo
lo::hotrestart: hot_restart register
_listener_fd: 127.0.0.1:8083, 10
2023-09-15T02:51:44.750135Z INFO vo
lo_thrift::server: [VOL0] server sta
rt at: Tcp(TcpListenerStream { inner
: PollEvented { io: Some(TcpListener
{ addr: 127.0.0.1:8083, fd: 10 }) })
})
xxx
propagate to 0 clients
COMMAND SAVED!!
xxx
propagate to 0 clients
COMMAND SAVED!!
xxx
propagate to 0 clients
COMMAND SAVED!!
[]

└─> mas1
└─> mas2
└─> sla1
└─> sla2
└─> proxy_srv
└─> proxy_client
```

测试样例

Graceful exit (10分)

收到退出信号时可以等待当前正在处理的连接执行结束之后再退出。

模拟处理服务过程

```
match _req.cmd {  
  RedisCommand::Ping => {  
    tokio::time::sleep(Duration::from_secs(10)).await;  
    if let Some(arg: Vec<FastStr>) = _req.args {
```

Ping后等待响应

```
vodis[connected]> ping
```

捕获Ctrl- C后进入等待

```
0.0.1:8080, socket: 324 }) } })  
2023-09-15T02:45:09.686504Z INFO volo_thrift::server: [VOLO] received s  
2023-09-15T02:45:09.686527Z INFO mini_redis: Ctrl-C received, shutting
```

等待服务处理完后退出

```
2023-09-15T02:54:00.345424Z INFO volo_thrift::server: [VOLO] server start at: Tcp(TcpListenerStream  
0.0.1:8080, socket: 336 }) } })  
2023-09-15T02:54:16.984492Z INFO volo_thrift::server: [VOLO] received signal, gracefully exiting no  
2023-09-15T02:54:16.984538Z INFO mini_redis: Ctrl-C received, shutting down  
2023-09-15T02:54:23.809971Z INFO mini_redis: New requests rejected whe shutting down.  
2023-09-15T02:54:23.809971Z INFO mini_redis: shutdown finally
```

Transactions (20分)

1. MULTI EXEC 命令实现。
2. WATCH 命令实现。

```
vodis[connected]> multi
[OK] M8Qpi
vodis[transaction]> watch 1
[OK] OK
vodis[transaction]> get 1
[OK] 1
vodis[transaction]> set 1 1
[OK] OK
vodis[transaction]> get 1 -t M8Qpi
[OK] OK
vodis[transaction]> set 2 -t M8Qpi
error: the following required arguments were not provided:
  <VALUE>

Usage: set --transaction-id <TRANSACTION_ID> <KEY> <VALUE> [TTYPE] [TNUM]

For more information, try '--help'.
vodis[transaction]> set 2 1 -t M8Qpi
[OK] OK
vodis[transaction]> exec
[OK] 1
[OK] OK
vodis[connected]> get 1
[OK] 1
vodis[connected]> get 2
[OK] 1
vodis[connected]> 
```

```
vodis[connected]> multi
[OK] c9GEc
vodis[transaction]> watch 1
[OK] OK
vodis[transaction]> set 1 1145
[OK] OK
vodis[transaction]> exec
2023-09-15T02:28:23.821340Z ERROR client_cli: Application(ApplicationError { kind: Appli
vodis[transaction]> 
```



其他实现细节

异步sync:

基于云函数rpc实现“远程interrupt”的效果;

布隆过滤器:

减少缓存击穿与读放大;

基于状态机的分支处理;

可扩展性:

低耦合, 可以简单从大作业要求修改到真实redis工作模型;

包装完善的命令行工具。

```
> ./target/debug/client-cli.exe
vredis[connected]> help
Usage: mini_redis <COMMAND>

Commands:
  ping      ping the server
  set       set a key-value pair
  get       get a value by key, nil if key not exist
  del       delete a key-value pair, error if key not exist
  exit      quit the client
  publish   publish a message to a channel
  subscribe subscribe a channel
  watch     watch a transaction
  multi     start a transaction
  exec      execute a transaction
  help      Print this message or the help of the given subcommand(s)

Options:
  -h, --help    Print help
  -V, --version Print version
vredis[connected]> |
```

```
> ./target/debug/server.exe -h
Usage: server.exe [OPTIONS] --ip <IP> --port <port>

Options:
  -n, --name <NAME>      Server name, used in default AOF file name. Will be randomly chosen if not provided
  -a, --aof <FILE>       Optional AOF file path
  -s, --slaveof <Master IP:PORT> Mark this Vredis instance as a slave
  -i, --ip <IP>
  -p, --port <port>
  -c, --cluster...       Sets a custom config file
                          --pre-run <PRE_RUN> Execute provided commands after initialization
  -h, --help             Print help
  -V, --version          Print version
```